Every Tongue

AI Bible Study for Every Language

AMD Developer Hackathon — Track 3 (Unicorn) Submission

The Mission

Multilingual AI Bible study spanning 120+ languages — from major world languages (Spanish, French, Arabic, Hindi, Portuguese) to low-resource and indigenous ones — built on AMD ROCm hardware, deployed and live. Hackathon demo used Swahili and Akuapem Twi to illustrate the full quality spectrum across the 7,000+ languages AI tools still underserve.

Submitted By

Bulletproof Automations
Johnathan Lightfoot
johnathan@bulletproofAutomations.com

Track

Track 3 — Unicorn

Model

Google Gemma 3 12B

Hardware

AMD Radeon gfx1100 · ROCm 7.2

Status

Live & Deployed

7,000 Languages. AI Serves a Handful.

The gap between what AI promises and what it delivers is not distributed equally. The communities most in need of AI productivity tools are precisely the ones current AI development ignores most completely.

The Gap Is Vast

Over 7,000 languages are spoken worldwide. Mainstream AI models perform well in fewer than 100. Languages spoken by the world's most isolated communities — from mid-resource regional languages to truly low-resource indigenous ones — are the ones AI ignores most completely. Not because the need is small, but because the training data is thin.

Ministry Workers Bear the Cost

Pastors, missionaries, and Bible translators in low-resource language communities must hand-translate every devotional, study guide, and discussion question — or go without entirely. There is no good AI tool for hundreds of languages spanning every continent, from Zarma to Bislama to Akuapem Twi and far beyond.

The Faith Sector Is Uniquely Affected

Bible translation is one of the largest, most sustained localization efforts on earth. Organizations like Wycliffe Bible Translators and SIL International operate across 2,000+ languages — yet even they lack AI-assisted drafting tools for the hardest cases. ScriptureFlow already supports 120+ languages across the full spectrum, from Spanish, French, Arabic, and Hindi down to genuinely low-resource languages.

The Irony

The communities that would benefit most from AI productivity tools are exactly the ones current AI development ignores. Every Tongue exists to close that gap — not by waiting for Big Tech to notice, but by building the infrastructure ourselves. The hackathon demo used Swahili and Akuapem Twi as a representative subset: one stronger-performing language, one genuinely low-resource one — illustrating the full quality spectrum our platform must handle.

The Technical Root Cause — and Our Solution

Why AI Fails Low-Resource Languages

Large language models learn from internet-scale text. Spanish and French have billions of web documents; Akuapem Twi has thousands. Models trained on this imbalance produce fluent output in major languages and unreliable, hallucinated text in low-resource ones — including **fabricated Scripture references**, a critical failure mode for ministry use where textual accuracy is doctrinal. Our hackathon demo used Swahili and Akuapem Twi to illustrate this full quality spectrum — from a strong-performing language to a genuinely low-resource one.

Our Core Differentiator

Every Tongue pulls real, expert-translated Scripture in both English and the target language from **ScriptureFlow**, our own multilingual Scripture data platform — covering 50+ languages spanning major world languages (Spanish, French, Arabic, Hindi, Portuguese) down to low-resource and indigenous ones. This parallel text is injected directly into the model prompt, so Gemma generates study materials anchored to **verified text** — not its own unreliable memory of what a verse might say.

1

Scripture Grounding

Parallel text from ScriptureFlow is injected into the prompt — verified, expert-translated source material anchors every generation across all 50+ supported languages.

2

Back-Translation Verification

Every AI draft is generated in both the target language and English simultaneously. Non-speaking supervisors can verify meaning survived the generation step — no native speaker required at every review stage.

3

Flagging for Human Review

The app automatically flags phrases below confidence thresholds and routes them for native-speaker review — making human oversight a structured part of the workflow, not an afterthought.

AMD Hardware: The Engine Behind Every Tongue

Every inference run in Every Tongue executes on real AMD hardware, configured end-to-end on the AMD Developer Cloud. The full evidence trail — ROCm configuration, vLLM setup, benchmark logs, and generation outputs — is committed to our public GitHub repository and verifiable by judges.

Technical Stack

Model: Google Gemma 3 12B — frontier open-weight multilingual model
Hardware: AMD Radeon GPU (gfx1100, RDNA3 architecture) on AMD Developer Cloud
Inference Engine: vLLM
Compute Platform: AMD ROCm 7.2 — AMD's open-source GPU compute platform enabling full PyTorch and vLLM compatibility on Radeon hardware

All AMD usage — ROCm config, vLLM setup, benchmark logs — committed to public repo and verifiable by judges.

AMD Radeon gfx1100 — Gemma 3 12B Throughput (ROCm 7.2 + vLLM)

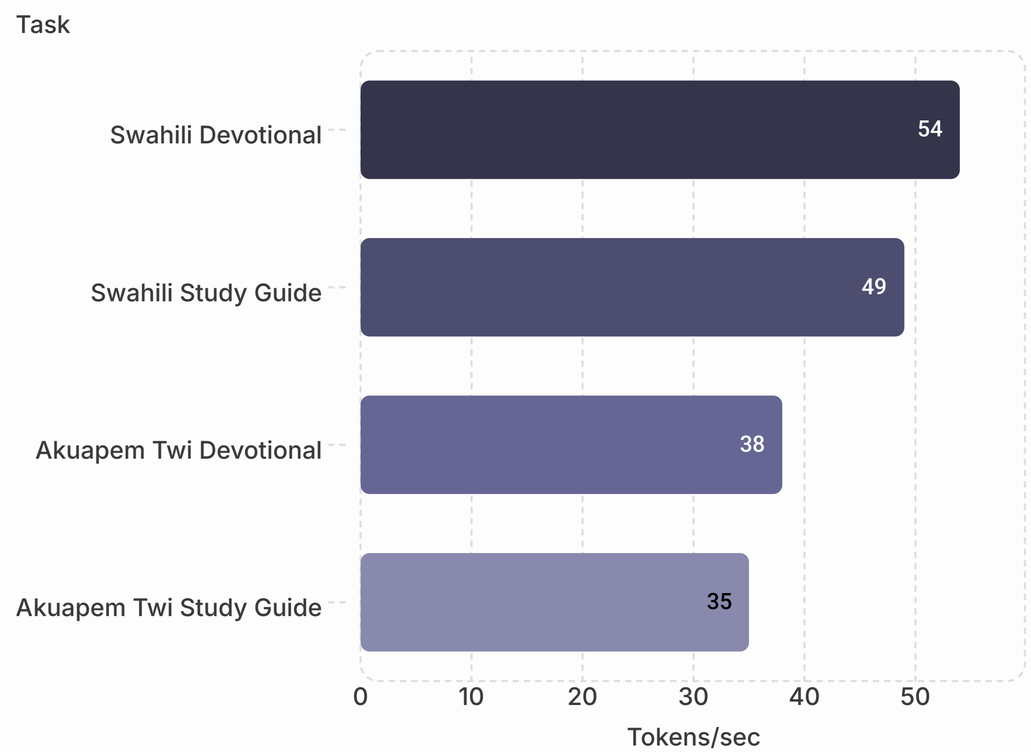

35–54 tokens/sec on real multilingual generation tasks, benchmarked and committed to the public GitHub repository. These two languages — Swahili (mid-resource) and Akuapem Twi (genuinely low-resource) — were selected as a representative subset from ScriptureFlow's 120+ supported languages spanning Spanish, French, Arabic, Hindi, Portuguese, and many others. Production-viable throughput on accessible AMD Radeon hardware — not enterprise-grade infrastructure.

Our Honest Thesis: The Gap IS the Product

Most submissions hide their weaknesses. We put ours on the cover slide. Here is what our testing actually shows — and why it strengthens, not undermines, the case for Every Tongue.

What Testing Revealed

Gemma 3 12B generates fluent, coherent Swahili Bible study content. For Akuapem Twi — a genuinely low-resource language — output quality degrades noticeably: grammar weakens, idiomatic expression suffers, and without grounding, Scripture references drift. This quality spectrum exists across all 50+ languages ScriptureFlow supports, from major world languages to truly low-resource ones.

We Don't Hide This — It's Our Argument

This performance gap is precisely why Every Tongue exists. If frontier AI already handled every low-resource language perfectly, there would be no problem to solve. **The gap proves the need.** Our hackathon demo uses Swahili and Akuapem Twi as representative points on that spectrum — not the full scope of the platform.

How We Bridge It

Scripture grounding (verified parallel text in the prompt) + mandatory human review flags + back-translation verification combine to make AI output genuinely usable even where the model is weakest — whether that's Spanish dialects, Arabic variants, Hindi scripts, or truly low-resource indigenous languages. We make the AI a capable first-draft tool, not a replacement for human expertise.

The Compounding Advantage

Every new language added to ScriptureFlow's parallel corpus improves grounding quality for that language — creating a compounding data advantage as the platform grows across its 120+ supported languages and beyond. The hard cases get easier over time, across the full spectrum from mid-resource to truly low-resource.

Demo strategy: We demo Swahili as the strong case — showing the system at its best. Akuapem Twi remains in the live app as the "hard case" demonstrating the real-world challenge and the value of our review workflow. These two languages represent the full quality spectrum ScriptureFlow navigates across 120+ languages. Judges see both ends of that spectrum.

What Every Tongue Generates — Live Demo

Content Types Generated

- Devotionals from a Scripture reference input
- Verse-by-verse study guides
- Discussion questions for small groups
- Passage summaries in target language + English back-translation

Languages Supported Today

- **Swahili** — strong case, production-quality output
- **Akuapem Twi** — hard case, showcased challenge with review workflow
- ScriptureFlow's **199 translations across 125 languages.** — from Spanish, French, Arabic, and Hindi to low-resource and indigenous languages worldwide

The Generation Workflow

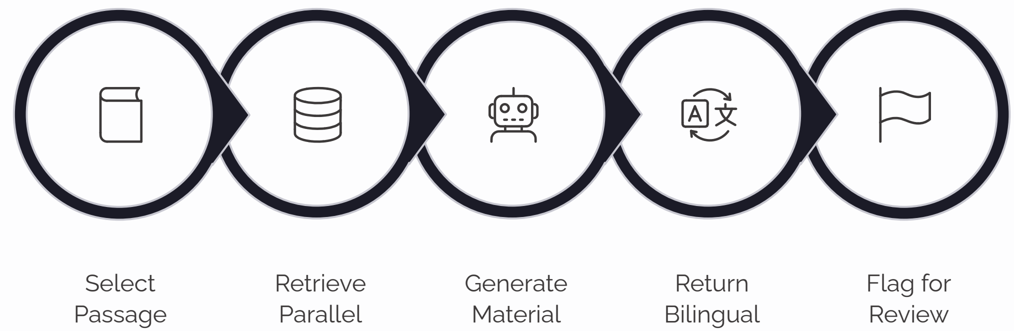

Live & Deployed

huggingface.co/spaces/exnav29/every-tongue — running now. Judges can interact directly.

Real AMD Inference

Working pipeline on AMD Radeon gfx1100 via ROCm 7.2 + vLLM — not a prototype or mockup.

Public Codebase

github.com/Exnav29/every-tongue — all AMD usage, benchmarks, and outputs committed and verifiable.

ScriptureFlow: The Platform Beneath the Product

Every Tongue is the first application built on ScriptureFlow — a proprietary multilingual Scripture data platform that constitutes the structural data moat beneath the entire product family.

What It Is

ScriptureFlow is a structured, queryable corpus of expert-translated Bible text spanning **120+ languages** — from major world languages (Spanish, French, Arabic, Hindi, Portuguese) down to low-resource and indigenous languages — purpose-built for AI grounding use cases. Accessible to AI applications via a clean API, it currently indexes **199 translations**. Expert-translated Scripture is the highest-quality parallel text available for many low-resource languages — it exists because of decades of investment by organizations like Wycliffe and SIL.

The Compounding Data Moat

Each new language added to ScriptureFlow improves grounding quality for that language across *all* products built on the platform — a structural advantage that grows with usage. The more languages we add, the better every product gets.

Platform Expansion Path

01

Every Tongue (Now)

Bible study generation for ministry workers across the full language spectrum — from major world languages to genuinely low-resource ones

02

Sermon & Translation Aids

Sermon preparation tools and oral Bible translation aids for field workers

03

Literacy Materials

NGO-facing content localization for education and literacy programs

04

Domain-Agnostic Expansion

Legal documents, health information, government content — the same infrastructure, a much larger market

Market Opportunity: Faith Sector and Beyond

The beachhead market is large, well-funded, and institutionally structured to adopt subscription software. The expansion path extends well beyond the faith sector.

Segment	Est. Users / Orgs	Model
Bible translation orgs (Wycliffe, SIL, affiliates)	2,000+ language programs	Org subscription
Missionaries & ministry workers	400,000+ globally	Free / individual tier
NGOs & international development	Thousands of orgs	Platform / API licensing
Government multilingual content	Emerging market	Enterprise licensing

Business Model

- **Free tier** for individual translators and missionaries — drives adoption and corpus feedback
- **Subscription tiers** for ministry organizations and NGOs — recurring revenue
- **Platform licensing** for ScriptureFlow API access — highest-margin, most scalable

Why Now

Open-weight models like Gemma 3, combined with AMD ROCm enabling affordable GPU inference on AMD Radeon hardware, make this economically viable for the first time. The infrastructure cost barrier has collapsed. The communities that need these tools most — spanning major world languages like Spanish, French, Arabic, Hindi, and Portuguese, down to mid-resource and truly low-resource languages — can now be served at a price point that works.

Why Every Tongue Wins on All Four Judging Criteria

The following maps our submission directly to each AMD hackathon judging criterion by name. Judges may score directly from this slide.

Creativity & Originality

Scripture-grounded AI generation for low-resource and mid-resource languages is a novel technique applied to a genuinely underserved problem. The back-translation verification loop and human-review flagging system are original workflow innovations. No comparable product exists in the market.

Product & Market Potential

A real platform (ScriptureFlow, 120+ languages spanning major world languages to low-resource and indigenous ones) beneath a real product (Every Tongue), targeting a large, well-funded, and genuinely underserved market. Clear expansion path from faith sector to NGO and government localization. Subscription business model with compounding data advantages.

Completeness

Live deployed application at huggingface.co/spaces/exnav29/every-tongue. Real AMD GPU inference running Gemma 3 12B via ROCm 7.2 and vLLM. Benchmarked throughput: 35–54 tokens/sec. Hackathon demo uses Swahili and Akuapem Twi as a representative subset — illustrating the full quality spectrum from strong-performing to genuinely low-resource languages. Public GitHub repository with all AMD evidence committed. **Judges can run it now.**

Meaningful AMD Compute Usage

Gemma 3 12B served on AMD Radeon GPU (gfx1100, RDNA3) via AMD ROCm 7.2 + vLLM on AMD Developer Cloud. Not a wrapper around a hosted API — actual GPU inference on AMD hardware, measured, documented, and reproducible. AMD ROCm is the reason this is affordable enough to serve every language AI underserves, from Spanish and Arabic to the world's most low-resource tongues.

The Ask: Fund the Languages AI Left Behind

> Every community deserves AI tools in their own language — and the languages that need it most are the ones we're building for first.

What We've Built

A working, deployed AI application serving real ministry workers across 50+ languages — from major world languages like Spanish, French, Arabic, and Hindi down to low-resource and indigenous languages — grounded in expert-translated Scripture, running on AMD ROCm hardware, with a live demo judges can use today. The hackathon demo focuses on Swahili and Akuapem Twi to illustrate the full quality spectrum, from strong-performing to genuinely low-resource languages. This is not a concept or wireframe. It is production infrastructure.

What Winning Enables

- Expand ScriptureFlow's language corpus beyond its current **199 translations across 125 languages.** toward the full spectrum of underserved world languages
- Build native-speaker reviewer network for quality assurance at scale
- Develop organization-tier subscription features for Wycliffe and SIL partnership pipeline
- Pursue formal corpus access and distribution partnerships with major translation organizations

Technical Stack Summary

Google Gemma 3 12B

Frontier open-weight multilingual model

AMD Radeon gfx1100 (RDNA3)

AMD Developer Cloud — accessible, affordable hardware

AMD ROCm 7.2 + vLLM

Full PyTorch & vLLM compatibility on Radeon hardware

ScriptureFlow Corpus

199 expert-translated multilingual Bible translations across 120+ languages

Live Demo: huggingface.co/spaces/exnav29/every-tongue

Public Repo: github.com/Exnav29/every-tongue

Contact & Resources

Every Tongue is live, open, and ready for judges to evaluate right now.

Team

Bulletproof Automations

Founder

Johnathan Lightfoot

Email

johnathan@bulletproofAutomations.com

Live Demo

The application is running now and accepts real inputs. Select a Scripture passage, choose Swahili or Akuapem Twi — a representative subset spanning strong-performing and genuinely low-resource languages — and observe grounded multilingual generation on AMD ROCm hardware.

huggingface.co/spaces/exnav29/every-tongue

Public Repository

All AMD usage evidence is committed and verifiable: ROCm configuration, vLLM setup, benchmark logs, and generation outputs. Full codebase is open for inspection.

github.com/Exnav29/every-tongue

AMD technical identifiers present throughout this submission as readable text: AMD · ROCm · Radeon · gfx1100 · vLLM · Gemma — all verified in the public repository.